

Build Model But we Handed Data To Be Balanced Using OverSampling

```
In [46]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier

# to ignore Errors
import warnings
warnings.filterwarnings("ignore", category=FutureWarning)
```

```
In [33]: # Read Data into Dataframe

df = pd.read_csv(r'C:\Users\dell\Desktop\powerbi projects\datasets\telecom\telecom_data.csv')
```

```
In [5]: # split the data into Features and Target

X = df.drop(['customerID', 'Churn'],axis=1)
y = df['Churn']
```

```
In [6]: # Show The Target Distribution Before OverSampling

y.value_counts()
```

Out[6]: No 5174
Yes 1869
Name: Churn, dtype: int64

- Apply over_sampling:-

```
In [17]: # over_sampling
from imblearn.over_sampling import RandomOverSampler

ros = RandomOverSampler()
X,y = ros.fit_resample(X,y)
```

```
In [18]: # Show The Target Distribution After OverSampling

y.value_counts()
```

Out[18]: No 5174
Yes 5174
Name: Churn, dtype: int64

Now Data Is Balanced

```
In [25]: # ordinal encoding

from sklearn.preprocessing import OrdinalEncoder

oe = OrdinalEncoder()
X = oe.fit_transform(X)
```

```
In [26]: # train test split

X_train,X_test,y_train,y_test=train_test_split(X,y,random_state=10)
```

- Build Model:-

```
In [15]: # RandomForestClassifier Model

from sklearn.ensemble import RandomForestClassifier

rfc = RandomForestClassifier()
rfc.fit(X_train,y_train)

y_pred = rfc.predict(X_test)
```

```
In [45]: # print The Evaluation Metrics

print("RandomForestClassifier Results")
print("-----")
print(classification_report(y_test,y_pred))
print("-----")
```

RandomForestClassifier Results					
	precision	recall	f1-score	support	
No	0.95	0.84	0.89	1318	
Yes	0.85	0.95	0.90	1269	
accuracy			0.90	2587	
macro avg	0.90	0.90	0.90	2587	
weighted avg	0.90	0.90	0.90	2587	

The Model precision = 95% with NO And 85% With Yes

The Model Recall = 85% with NO And 95% With Yes

The Accuracy is 90%

Done